

I have built my career around the principle that AI only creates lasting value when it is built on secure, governed and well-integrated foundations. AI cannot operate effectively as an isolated capability, it depends on trusted data, scalable platforms, strong integration and clear operational controls. This is what draws me to AI and Platform Technical Architect roles within the Civil Service, where I want to help teams adopt AI safely and at scale.

My background combines Azure cloud architecture with deep AI enablement experience, underpinned by Microsoft certifications as an Azure AI Engineer Associate and Azure Solutions Architect Expert. I have designed secure foundations for Azure OpenAI, Microsoft Foundry and the wider Azure AI Services estate, including Azure AI Search, Vision, Speech, Language, Translator, Document Intelligence, Content Safety and Cognitive Services, and I have architected multi-agent systems using Microsoft Agent Framework, Semantic Kernel, LangChain, LangGraph, LangSmith, AutoGen, CrewAI and MCP (Model Context Protocol) for standardised tool and context access across agents.

I treat AI architecture as part of the wider enterprise environment rather than an experimental capability sitting outside established controls. I designed patterns for RAG pipelines, vector search and agent memory using private endpoints, managed identities, restricted network access, centralised monitoring and Azure API Management, so that AI services inherited the same security posture as the rest of the estate. I also embedded guardrails, content safety controls, evaluation frameworks and model routing so that AI outputs remained safe, explainable and cost-effective across multiple models, drawing on Hugging Face models, fine-tuning and prompt engineering to adapt capabilities to specific use cases. These controls gave teams a governed route to AI and agentic capabilities while supporting authentication, throttling, monitoring and cost management.

I integrated AI capabilities with core Azure application, data and messaging services, including Functions, App Service, Service Bus, Event Grid, Cosmos DB, Azure SQL and Storage Account. I considered how information moved through each solution, where access should be restricted, how components would recover from failure and how service consumption could be monitored, so that AI and agentic workflows could be adopted through reusable enterprise patterns without weakening security or operational governance.

I documented significant AI architecture decisions, model choices and trade-offs through Architecture Decision Records, giving transparency around why decisions were made, reducing reliance on knowledge held by individuals, and helping engineers understand, operate and evolve AI systems responsibly over time.

Architecture also requires communication, influence and collaboration. I have worked with product, engineering, cyber security, data and operations stakeholders to evaluate AI use cases, explain risks around bias, safety and cost, and agree practical, proportionate controls. I adapt technical detail for different audiences, relating AI architecture decisions to service reliability, explainability, security exposure and operating cost, and where teams were

concerned that governance would slow adoption, I demonstrated how reusable, governed patterns could accelerate safe delivery rather than obstruct it.

The approach increased AI delivery efficiency by 40%, reduced cloud operating costs by 30%, and cut critical security issues by 65% across the wider platform. It created a governed, reusable AI capability that allowed teams to adopt agentic AI with embedded security, evaluation and operational controls, supporting reliable and responsible public-sector service delivery.

I am applying to the Civil Service because I want to help ensure that AI adopted across public services is safe, explainable, and genuinely useful rather than experimental. I would bring hands-on Azure and agentic AI expertise, enterprise architectural judgement and a practical, risk-based approach to governance. I measure effective AI architecture by whether teams can understand it, operate it safely and trust its outputs, and I am motivated by the opportunity to help the Civil Service build trusted, secure and cost-effective AI capabilities that support better public services.