My current role sits between technical architecture, platform engineering and AI enablement. I remain close enough to delivery to understand how architectural decisions work in practice, while also providing the standards, governance and technical direction needed for teams adopting AI. Some days this means reviewing an agent or RAG design; on others, it means documenting an architectural decision, resolving safety concerns or helping a team adopt an approved AI pattern.

DAY-TO-DAY ENGINEERING

I design and govern the shared Azure AI platform foundations used by multiple product teams to build and operate agentic and generative AI services, backed by my certification as an Azure AI Engineer Associate. My day-to-day responsibilities include reviewing AI solution designs, assessing use cases and selecting appropriate models, frameworks and Azure AI services based on safety, scalability, latency, explainability and cost.

I develop reference architectures for multi-agent systems using Microsoft Foundry, Microsoft Agent Framework, Semantic Kernel, LangChain, LangGraph, LangSmith, AutoGen, CrewAI and MCP (Model Context Protocol), and for the wider Azure AI Services estate, including Azure AI Search, Vision, Speech, Language, Translator, Document Intelligence, Content Safety and Cognitive Services. I define patterns for RAG pipelines, vector search, agent memory and function-calling tools.

I remain involved throughout delivery, resolving design issues, reviewing implementations and improving evaluation frameworks, guardrails and model-routing configurations based on team feedback, drawing on Hugging Face models, fine-tuning and prompt engineering where needed.

TECHNICAL LEADERSHIP

I provide technical direction for AI architecture and secure AI adoption. I establish architecture principles, reusable agent and RAG patterns and engineering standards, then translate them into automated components that teams can apply consistently.

I ensure that AI services are provisioned through version-controlled deployments and that safety and security validation is completed before release. Azure AI Content Safety, guardrails, evaluation frameworks and automated compliance checks are embedded into delivery workflows.

I document significant options, risks and trade-offs through Architecture Decision Records. I also define secure patterns for Azure OpenAI and Azure AI Services using managed identities, private connectivity, controlled API access, monitoring and cost-management controls.

COLLABORATION AND STAKEHOLDER WORK

I work with product, engineering, cyber security, data and operations stakeholders to understand AI use cases, evaluate technical options and agree practical, proportionate controls. I communicate architecture decisions in terms relevant to each audience, including reliability, explainability, delivery timescales, security exposure and cost. Where teams are concerned that governance may slow AI adoption, I demonstrate how reusable, governed patterns can reduce manual effort and accelerate safe delivery. I support teams through implementation and use their experience to improve platform standards and reference architectures.

GROWING THE TEAM

I help engineers and architects understand not only how to apply an AI pattern, but also why particular decisions were made — around model choice, safety controls or evaluation design. I share guidance through design reviews, reference architectures, Architecture Decision Records and practical implementation support.
My objective is to reduce dependency on individual specialists by creating reusable knowledge and enabling teams to adopt AI more independently and safely. This approach has contributed to an 85% reduction in provisioning time, a 90% improvement in deployment consistency, a 65% reduction in critical security findings, 40% faster delivery cycles and a 30% reduction in cloud operating costs.